

Document Similarity Benchmark

Jan-Gabriel Mylius

Lehrstuhl für Datenbanksysteme,
Institut für Informatik,
Ruprecht-Karls-Universität Heidelberg
mylius@posteo.de

10. February 2020, Heidelberg

Motivation

- ▶ NLP is often interested in meaning of a sentence
- ▶ Comparing the semantic similarity is a common task

Sentence 1	Sentence 2
The man is operating a stenograph.	The man is typing on a machine used for stenography.
"Fairies don't exist" – fine.	"Satyrs don't exist" – fine.

The Architecture

The benchmark suite consists of two essential elements:

- ▶ the algorithms
- ▶ the datasets

We want a framework:

- ▶ That compares algorithms on different datasets
- ▶ Is easily expandable
- ▶ Provides commonly used algorithms
- ▶ Provides common datasets

Algorithms

There is a template class for implementing custom algorithms. All algorithms do the following 3 things

- ▶ train - Training sentences pairs and scores available
- ▶ encode sentences - convert sentence to a vector
- ▶ compare sentences - two sentences given, number should be returned

Preimplemented Algorithms

For the benchmark suite I implemented the following:

- ▶ Bag of Words
- ▶ Bag of Words with lemmatization
- ▶ Word2Vec[Mik+15]
- ▶ BERT[Dev+18]
- ▶ Word Mover's Distance[Kus+15]
- ▶ Doc2Vec[LM14]


Functional Principle

How do these measures function:

1. Convert words to vector
2. Compare vectors using a similarity metric

For some algorithms several similarity metrics were implemented:

- ▶ Cosine similarity
- ▶ Euclidean similarity
- ▶ Jaccard similarity



Euclidean vs. Cosine Distance

The dataset class

The dataset class has the following essential functions:


- ▶ loads a dataset
 - ▶ STS
 - ▶ SICK
 - ▶ custom
- ▶ runs an algorithm
- ▶ norms results to match the datasets scale
- ▶ stores results for all algorithms that have run

The benchmark


What the benchmark does:

- ▶ Accepts command line parameters:
 - ▶ No parameters: Run all algorithms
 - ▶ Else: Run given algorithms
- ▶ Measures:
 - ▶ Training Time
 - ▶ Run Time
 - ▶ Pearson correlation
 - ▶ Spearman's correlation
 - ▶ Mean squared error
- ▶ Outputs results as json file


Overview of all results



SICK – results without l2 and jaccard



STS – results without l2 and jaccard



Discussion of Results

- ▶ l2, jaccard and wmd similarity not useful
- ▶ Filtering stop words improves results
- ▶ Lemmatization improves results
- ▶ Doc2Vec needs more training data
 - ▶ Maybe allow additional training data in future version
 - ▶ There is no currently compatible pretrained Doc2Vec model
 - Some one should train one for current gensim
 - ▶ More and larger datasets would be interesting to have
- ▶ BERT and Word2Vec not reliably better than BoW with lemmatization

The Datasets – Lessons Learned

Set	Sentence 1	Sentence 2	Score
SICK	The man is operating a stenograph.	The man is typing on a machine used for stenography.	4.2(1-5)
STS	"Fairies don't exist" – fine.	"Satyrs don't exist" – fine.	1.2(0-5)
SICK	A man is jumping into an empty pool	There is no biker jumping in the air	1.6(1-5)

- ▶ Similarity and relatedness are not identical
 - ▶ Semantic similarity/relatedness are hard to quantify
 - ▶ STS/SICK ratings don't match each other
 - inter annotator consistency?
 - ▶ Finding consens on scale difficult
 - annotators training
- different annotation method might be better

Interface

Reference sentence:

Die Koalition droht am Streit über die Grundrente mit oder ohne Bedürftigkeitsprüfung zu zerbrechen. Während Unions-Fraktionsvize Linnemann ein „neues Aufbruchssignal“ fordert, ruft der CSU-Landesgruppenchef zum Kompromiss auf.

Sentence 1:


Nach Ansicht von Union-Fraktionsvize Carsten Linnemann sollte die Bundesregierung nicht an der Diskussion über die Grundrente scheitern. „Ich würde die große Koalition an dieser Frage nicht platzen lassen“, sagte der CDU-Politiker am Montag im ZDF-„Morgenmagazin“. Er würde seine Überzeugung aber nicht über Bord werfen, aus Rücksicht vor einer Mitgliederentscheidung bei der SPD.

Sentence 2:

Wegen noch offener Fragen bei der Grundrente war das für Montagabend geplante Spitzentreffen der Koalition auf den 10. November verschoben worden. Es gebe noch offene Punkte, die im Laufe dieser Woche sorgfältig geklärt werden sollten, teilte die CDU mit. Von der SPD hieß es, die Verschiebung sei von der Union ausgegangen. Die Arbeitsgruppe zu dem Thema habe gute Vorarbeit geleistet.

Which sentence is more similar to the reference sentence?

- ▶ Annotations are dictionaries of tuples
- ▶ Allows creation of different task format
- ▶ Annotations can be viewed as graph for each reference
 - ▶ Nodes: sentences
 - ▶ Edges: more similar relationship
 - ▶ Edge weights: number of annotators agreeing



Idea: Page-Rank-Algorithm to create scores

Conclusion

- ▶ Lemmatization and stop word filtering work
- ▶ Cosine similarity metric of choice
- ▶ Difference between similarity and relatedness matters

Future work:

- ▶ Create option to generate additional training data
- ▶ Implement cutting edge algorithms
- ▶ Implement Page-Rank for annotator
- ▶ Create dataset to test similarity vs relatedness

Bibliography I

- [Dev+18] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. “**Bert: Pre-training of deep bidirectional transformers for language understanding**”. In: *arXiv preprint arXiv:1810.04805* (2018).
- [Kus+15] Matt Kusner, Yu Sun, Nicholas Kolkin, and Kilian Weinberger. “**From word embeddings to document distances**”. In: *International conference on machine learning*. 2015, pp. 957–966.
- [LM14] Quoc Le and Tomas Mikolov. “**Distributed representations of sentences and documents**”. In: *International conference on machine learning*. 2014, pp. 1188–1196.
- [Mik+15] Tomas Mikolov, Kai Chen, Gregory S Corrado, and Jeffrey A Dean. *Computing numeric representations of words in a high-dimensional space*. US Patent 9,037,464. 2015.
- [Pag+99] Lawrence Page, Sergey Brin, Rajeev Motwani, and Terry Winograd. *The pagerank citation ranking: Bringing order to the web*. Tech. rep. Stanford InfoLab, 1999.

Outline

Motivation

- Motivation

Implementation

- The Architecture

Results

- Results

- About Dataset

The Annotator


- Annotator – Interface

Conclusion

- Conclusion

Bibliography

The linear case




Expected result from a rational single user:

- ▶ Rank sentences, then divide by number of sentences \approx score
- ▶ Problems:
 - ▶ Users aren't rational
 - ▶ Users might disagree

What does a realistic result look like?

The realistic case



Problems:


- ▶ Circular relationships exist
- ▶ Ranking is non obvious

Naive approach:

- ▶ Add up the edge weights for each node:
 - ▶ Node 3: 16, Node 5: 7
 - ▶ Does this match our expectations?

Better Solution?

The realistic case



Problems:

- ▶ Circular relationships exist
- ▶ Ranking is non obvious


Naive approach:

- ▶ Add up the edge weights for each node:
 - ▶ Node 3: 16, Node 5: 7
 - ▶ Does this match our expectations?

Better Solution?

Page-Rank Algorithm[Pag+99]

The realistic case



Problems:

- ▶ Circular relationships exist
- ▶ Ranking is non obvious

Naive approach:

- ▶ Add up the edge weights for each node:
 - ▶ Node 3: 16, Node 5: 7
 - ▶ Does this match our expectations?

Better Solution?

Page-Rank Algorithm[Pag+99]

Suggested Page-Rank Implementation

- ▶ make sure reference sentence is rated to itself
- ▶ run page-rank over annotations
- ▶ norm the results over all sentences

Reference Sentence compared to itself \longrightarrow upper bound.
Therefore this should give us decent scores.

Advantages:

- ▶ Easier to annotate than 1-5 scale.
- ▶ Quicker to annotate.
- ▶ Direct comparison data exists for sentence pairs.